

# SARX Feature Overview

Luke Miller

November 15, 2025

## SARX Feature Overview

- **Substrate-Inspired ARX Stream Cipher (SARX):**

- 256-bit ARX-256 keystream in CTR mode: 32-byte blocks, 8-round column+diagonal ARX core with feedforward, implemented as a stateless keystream generator.
- Per-vault keying: the encryption key `base_key32` is derived as  $\text{BLAKE3}(\text{password} \parallel \text{timestamp\_ns})$ , binding each keystream to a specific password + wall-clock timestamp so streams are never reused across vaults.
- ARX core selected by search: the round structure and constants were chosen via an ARX parameter search tuned for near-ideal avalanche and near-isometry in the bit metric, rather than via any Cantorian or lattice construction.
- Unicode-aware keying: accepts any UTF-8 password (any language, emoji, symbols); policy enforces 30–100 Unicode codepoints with no normalization or downcasting, so no entropy is silently discarded.
- Pure ARX keystream: the legacy 2D postmix layer is disabled in v1.1.0; the shipping cipher is an “honest” ARX-256 stream with a well-defined, auditable core.
- Full AEAD-style vaults: encryption is stream XOR with the ARX keystream, and every vault carries a keyed BLAKE3 MAC over the header (as associated data) and ciphertext; any wrong password or modification is detected before decryption.
- Parallelizable: the keystream API is byte-seekable via offsets; the CLI slices files into multi-MiB chunks and encrypts/decrypts them in parallel with Rayon, reaching multi-GB/s throughput on desktop CPUs.
- Cross-platform: Rust core plus C mirror; builds on Linux, macOS, and Windows, and the ARX core is suitable for ARM, x86\_64, RISC-V, MIPS, PowerPC, and WebAssembly.

- **Security and Randomness:**

- Statistical quality: the ARX-256 stream passes NIST SP800-22, Dieharder, and PractRand batteries in long-run testing; output is empirically indistinguishable from random under standard suites.
- Strict avalanche behavior: single-bit flips in the key or plaintext propagate through the ARX core to flip  $\approx 50\%$  of output bits (about 32 of 64 per word) in experiments, matching the ideal avalanche target.
- Balanced Hamming weight: keystream words exhibit stable Hamming weight distribution over long runs; there is no observable bias drift or positional structure in the bit metric.

- **Password, KDF, and Key Management:**

- Primary KDF: Argon2id with parameters  $t = 3$ ,  $m = 2^{17}$  KiB (128 MiB),  $lanes = 1$  derives a variable-length stream key  $k_{\text{stream}}$  plus a fixed 256-bit MAC key  $k_{\text{mac}}$  from the UTF-8 password and a 256-bit salt.
- Thermodynamic hardening (**-thermo**): an optional mode that runs a large ARX-based random walk over a 512 MiB scratch buffer, then folds a BLAKE3 digest of the walk back into the KDF output to raise the per-guess energy and memory cost of brute force.
- Physically-informed parameters: KDF settings and thermo-hardening are calibrated against Landauer and Bekenstein bounds via dedicated bit-energy and brute-force cost experiments, tying the effective attack cost to a concrete bit-energy floor.
- Header-bound KDF: each vault header records the salt, Argon2 parameters, and a KDF identifier (Argon2id vs. Argon2id+thermo), so keys are reproducible for decryption but distinct across vaults.
- Native support for USB “sigilbook” tokens for password storage and retrieval, with atomically scoped export: passwords never hit disk outside the hardware token.
- CLI and GUI workflows for encrypting, decrypting, verifying, and managing vaults, with enforced 30–100 Unicode codepoint password policy.

- **Performance and Edge/IoT Integration:**

- Desktop performance: benchmarks of the ARX-256 keystream show keystream generation above 1 GB/s and end-to-end encryption in the hundreds of MB/s to multi-GB/s range on commodity multi-core CPUs, depending on thread count and KDF settings.
- Streaming-friendly: the ARX-256 keystream is seekable via a byte offset and supports long-lived streams without re-keying, ideal for pipe-style encryption and large files.
- Minimal runtime state: aside from KDF memory, the cipher core keeps only a 256-bit key and counter; memory overhead for encrypt/decrypt is essentially the chunk buffer.
- Live cam/mic encryption: supports real-time audio/video vaulting on commodity hardware, suitable for privacy-preserving surveillance and recording; mobile and SBC performance is being re-benchmarked against the new core.

- **Usability and Developer Features:**

- Modern desktop GUI (Rust, `egui`): file manager, drag-and-drop vault creation, batch tools, and sigilbook USB integration.
- CLI dispatcher with subcommands for encryption, decryption, verification, benchmarking, and optional **-thermo** mode to test hardened KDF settings.
- Structured, human-readable vault format: versioned headers encoding salt, timestamp, nonce, Argon2 parameters, KDF ID, plus a 32-byte keyed BLAKE3 MAC tag per vault.
- Modular, extensible core: the ARX keystream, KDF, and MAC layers are cleanly separated, making it easy to embed SARX into new frontends and hardware targets.
- UX niceties: when the sigilbook USB is present, vaults can be auto-decrypted and re-encrypted on double-click without exposing passwords or derived keys to disk.